

Just tired of endless loops!
or parallel: Stata module for parallel computing

George G. Vega Yon¹ Brian Quistorff²

¹University of Southern California
vegayon@usc.edu

²Microsoft AI and Research
Brian.Quistorff@microsoft.com

Stata Conference Baltimore
July 27–28, 2017

Thanks to Stata users worldwide for their valuable contributions. The usual disclaimers applies.

Agenda

Motivation

What is it and how does it work

Benchmarks

Syntax and Usage

Concluding Remarks

Motivation

- ▶ Both computation power and size of data are ever increasing
- ▶ Often our work is easily broken down into independent chunks
- ▶ Implementing parallel computing, even for these “embarrassingly parallel” problems, however, is not easy.
- ▶ Stata/MP exists, but only parallelizes a limited set of internal commands, not user commands.
- ▶ `parallel` aims to make this more convenient.

Motivation

What is it and how does it work

Benchmarks

Syntax and Usage

Concluding Remarks

What is it and how does it work

What is it?

- ▶ Inspired by the R package “snow” (several other examples exists: HTCondor, Matlab’s Parallel Toolbox, etc.)
- ▶ Launches “child” batch-mode Stata processes across multiple processors (e.g. simultaneous multi-threading, multiple cores, sockets, cluster nodes).
- ▶ Depending on the task, can reach near linear speedups proportional to the number of processors.
 - ▶ Thus having a quad-core computer can lead to a 400% speedup.

Simple usage

Serial:

- ▶ `gen v2 = v*v`
- ▶ `do byobs_calc.do`
- ▶ `bs, reps(5000): reg price foreign rep`

Parallel:

- ▶ `parallel: gen v2 = v*v`
- ▶ `parallel do byobs_calc.do`
- ▶ `parallel bs, reps(5000): reg price foreign rep`


What is it and how does it work

How does it work?

- ▶ Method is *split-apply-combine* like MapReduce.

What is it and how does it work

How does it work?



What is it and how does it work

How does it work?

- ▶ Method is *split-apply-combine* like MapReduce. Very flexible!
- ▶ Straightforward usage when there is observation- or group-level work
- ▶ If each iteration needs the entire dataset, then use procedure to split the tasks and load the data separately. Examples:
 - ▶ Table of seeds for each bootstrap resampling
 - ▶ Table of parameter values for simulations
- ▶ If the list of tasks is data-dependent then the “nodata” alternative mechanism allows for more flexibility.

Implementation

Some details

- ▶ Uses shell on Linux/MacOS. On Windows we have a compiled plugging allowing:
 - ▶ Functionality when the parent Stata is in batch-mode
 - ▶ Seamless user experience by launching the child programs in a hidden desktop (otherwise GUI for each steals focus)
- ▶ For a Linux/MacOS cluster with a shared filesystem (e.g. NFS) and ssh-like commands, can distribute across nodes.
 - ▶ New feature so we'd appreciate help from the community to extend to other cluster settings (e.g. PBS)
- ▶ Make sure that child tempnames or tempvars don't clash with those coming from parent.
- ▶ Passes through programs, macros and mata objects, but NOT Stata matrices or scalars. No state but datasets are returned to parent.
- ▶ Recover gracefully from child failures. Currently no re-try support.

Motivation

What is it and how does it work

Benchmarks

Syntax and Usage

Concluding Remarks

Benchmarks

Bootstrap with parallel bs

```
sysuse auto, clear expand 10

// Serial fashion
bs, rep($size) nodots: regress mpg weight gear foreign

// Parallel fashion
parallel setclusters $number_of_clusters
parallel bs, rep($size) nodots: regress mpg weight gear foreign
```

Problem size	Serial	2 Clusters	4 Clusters
1,000	2.93s ×2.69	1.62s ×1.48	1.09s ×1.00
2,000	5.80s ×2.85	3.13s ×1.54	2.03s ×1.00
4,000	11.59s ×3.01	6.27s ×1.62	3.86s ×1.00

Table: Absolute and relative computing times for each run of a basic bootstrap problem. For each given problem size, the first row shows the time in seconds, and the second row shows the relative time each method took to complete the task relative to using parallel with four clusters. Each cell represents a 1,000 runs.

Benchmarks

Simulations with `parallel sim`

```
prog def mysim, rclass
  // Data generating process
  drop _all
  set obs 1000
  gen eps = rnormal()
  gen X = rnormal()
  gen Y = X*2 + eps

  // Estimation
  reg Y X
  mat def ans = e(b)
  return scalar beta = ans[1,1]
end

// Serial fashion
simulate beta=r(beta), reps($size) nodots: mysim

// Parallel fashion
parallel setclusters $number_of_clusters
parallel sim, reps($size) expr(beta=r(beta)) nodots: mysim
```

Benchmarks

Simulations with `parallel sim` (cont.)

Problem size	Serial	2 Clusters	4 Clusters
1000	2.19s ×3.01	1.18s ×1.62	0.73s ×1.00
2000	4.36s ×3.29	2.29s ×1.73	1.33s ×1.00
4000	8.69s ×3.40	4.53s ×1.77	2.55s ×1.00

Table: Absolute and relative computing times for each run of a simple Monte Carlo exercise. For each given problem size, the first row shows the time in seconds, and the second row shows the relative time each method took to complete the task relative to using parallel with four clusters. Each cell represents a 1,000 runs.

Code for replicating this is available at
<https://github.com/gvegayon/parallel>

Motivation

What is it and how does it work

Benchmarks

Syntax and Usage

Concluding Remarks

Syntax and Usage

Setup

```
parallel setclusters #|default [, force hostnames(namelist)]
```

Main command types

```
parallel [, by(varlist) programs(namelist) mata seeds(string) randtype(random.org|datetime)  
nodata]: stata_cmd
```

```
parallel do filename [, by(varlist) programs(namelist) mata seeds(string)  
randtype(random.org|datetime) nodata]
```

Helper commands

```
parallel bs [, expression(exp_list) programs(namelist) mata seeds(string)  
randtype(random.org|datetime) bs_options]: stata_cmd
```

```
parallel sim [, expression(exp_list) programs(namelist) mata seeds(string)  
randtype(random.org|datetime) sim_options)]: stata_cmd
```

```
parallel append [files], do(command/dofile) [in(in) if(if) expression(expand_exp)  
programs(namelist) mata seeds(string) randtype(random.org|datetime)]
```

Additional Utilities

```
parallel version/clean/printlog/viewlog/numprocessors
```


Debugging

- ▶ Use **parallel printlog/viewlog** to view the log of the child process (includes some setup code as well). Can **set trace** in the child do-file or command to see more.
- ▶ Auxiliary files created during process (harder to use):
 - ▶ (Unix) `__pllID_shell.sh`
 - ▶ `__pllID_dataset.dta`
 - ▶ `__pllID_doNUM.do`
 - ▶ `__pllID_glob.do`
 - ▶ `__pllID_dtaNUM.dta`
 - ▶ `__pllID_finitoNUM`
- ▶ Can keep these around by specifying the **keep** or **keeplast** options

Syntax and Usage

Recommendations on its usage

parallel suits ...

- ▶ Repeated simulation
- ▶ Extensive nested control flow (loops, while, ifs, etc.)
- ▶ Bootstrapping/Jackknife
- ▶ Multiple MCMC chains to test for convergence (Gelman-Rubin test)

parallel doesn't suit ...

- ▶ (already) fast commands
- ▶ Regressions, ARIMA, etc.
- ▶ Linear Algebra
- ▶ Whatever Stata/MP does better (on single machine)

Use in other Stata modules

- ▶ EVENTSTUDY2: Perform event studies with complex test statistics
- ▶ MIPARALLEL: Perform parallel estimation for multiple imputed datasets
- ▶ Synth_Runner: Performs multiple Synthetic Control estimations for permutation testing

Concluding Remarks

- ▶ Brings parallel computing to many more commands than Stata/MP
- ▶ Its major strengths/advantages are in simulation models and non-vectorized operations such as control-flow statements.
- ▶ Depending on the proportion of the algorithm that can be parallelized, it is possible to reach near to linear scale speedups.
- ▶ We welcome other user commands optionally utilizing `parallel` for increased performance.
- ▶ Install, contribute, find help, and report bugs at <http://github.com/gvegayon/parallel>

Thank you very much!

George G. Vega Yon¹ Brian Quistorff²

¹University of Southern California
vegayon@usc.edu

²Microsoft AI and Research
Brian.Quistorff@microsoft.com

Stata Conference Baltimore
July 27–28, 2017